

Research-informed conceptual essay

Are we afraid of AI, or of ourselves?

AI as amplifier, moral buffer and moral alibi

Feliks Mamczur
ORCID: 0009-0001-0715-0517

Publication	Prompted Psyche	Published	July 4, 2026
Revised	July 13, 2026	Version	2.1
Language	English	License	CC BY 4.0
DOI	10.5281/zenodo.21340181	Canonical article	promptedpsyche.com

Suggested citation

Mamczur, F. (2026, July 4). Are we afraid of AI, or of ourselves? AI as amplifier, moral buffer and moral alibi (Version 2.1). Prompted Psyche. <https://doi.org/10.5281/zenodo.21340181>

This is the publication copy associated with the reserved Zenodo DOI. The work is a narrative synthesis, not a systematic review, not original empirical research and not peer-reviewed.

Publication map

Contents

- 1** Abstract
- 2** Scope and method
- 3** Evidence profile
- 4** Why the evil-AI story is attractive
- 5** Causation is not moral responsibility
- 6** Moral disengagement, delegation and distance
- 7** AI as amplifier, moral buffer and moral alibi
- 8** What experiments actually show
- 9** Institutions make the alibi credible
- 10** The opposite danger: gaps and crumple zones
- 11** Hypothetical case: the retention score
- 12** A responsibility-tracing protocol
- 13** Implications for design and governance
- 14** Limits of the argument
- 15** Conclusion
- 16** Sources and further reading

A recruitment team is reviewing a shortlist prepared with an AI-supported scoring system. One candidate has disappeared from the final group. A manager asks why. No one points to a single disqualifying fact. Instead, someone says: "The system marked the application as high risk."

The sentence sounds descriptive, almost passive. Yet the system did not decide what risk should mean. It did not choose which historical records counted as evidence, how much uncertainty was acceptable, who would be excluded or whether a human reviewer had enough time to challenge the score. People and institutions made those choices, even if no one person made all of them. This is a problem of distributed human agency, not machine intention.

Once harm occurs, the same sentence can become an explanation: the system indicated it. It can also become a shield. A chain of human decisions is compressed into an event that appears to have happened inside the machine.

This is one reason the image of a dangerous, autonomous AI is so compelling. It gives agency a visible address. The machine becomes the actor, while the less dramatic work of setting goals, buying systems, defining metrics, approving deployment and accepting outputs fades into the background.

AI can influence an outcome without becoming its moral author. When "the system decided" replaces a traceable account of goals, control and approval, technology can begin to function as a moral alibi.

Abstract

This research-informed conceptual essay asks whether stories about independently dangerous AI obscure the human intentions, incentives and organizational choices that give AI systems practical force. It is a narrative synthesis, not a systematic review or original empirical study. It brings together moral psychology, research on decision delegation and mind attribution, human factors and sociotechnical accounts of responsibility.

The central argument is that current AI systems can contribute causally to harmful outcomes without being moral agents in the human sense. They can extend a goal's reach, increase the distance between choice and consequence, and support a credible account in which a person or institution presents an outcome as the machine's decision. I describe these roles as AI as amplifier, AI as moral buffer and AI as moral alibi. They are interpretive categories for this essay, not validated psychological constructs, a diagnostic taxonomy or a measurement scale.

Experimental evidence gives the argument a bounded empirical basis. In specific moral, organizational and behavioral tasks, responsibility shifting or avoidance helped explain why participants delegated difficult choices to AI. Indirect machine delegation increased dishonest behavior in the tested tasks, while observers attributed blame or intention-like mental states to artificial agents. These findings do not show that every delegation is evasive, that users always anthropomorphize systems or that AI deserves moral blame.

Yet opacity and distributed control create a different danger. Learning systems may produce outcomes that are difficult to foresee, while authority is spread across designers, vendors, managers and operators. Responsibility may genuinely be hard to assign. A low-level human may also become a "moral crumple zone," absorbing blame for a system they had little power to change. A defensible account must distinguish causal contribution, moral responsibility, accountability, blame and legal liability. It must ask who defined the goal, knew the limits, could

intervene, approved the action and remains responsible for explanation and repair.

Scope and method

The evidence was selected for conceptual relevance rather than through a preregistered search strategy. Priority went to peer-reviewed experiments, foundational psychological theory, human-factors research, sociotechnical analysis and official risk-management guidance. The final set contains 28 sources: 26 peer-reviewed journal or conference publications and two reports, with no preprints.

The sources ask different questions. Vignette studies measure how observers distribute blame. Behavioral experiments test delegation under incentives. Moral psychology describes processes that can loosen the link between standards and action. Human-factors research examines what happens when operators retain formal responsibility but lose information, skill or effective control. Philosophy and science and technology studies clarify concepts such as responsibility gaps, agency laundering and meaningful human control.

These literatures cannot be combined into one effect size. A study of an autonomous vehicle is not a study of a language model, and a laboratory task is not an organization. I use the findings to constrain a conceptual argument, not to estimate how often AI becomes an alibi in the world. Claims about law are outside the article's scope. Legal liability depends on facts, jurisdiction and applicable rules.

Evidence profile

The essay draws on 28 sources selected for conceptual relevance. The classification below separates empirical evidence from reviews, philosophical analysis and institutional guidance. It does not turn a narrative synthesis into a systematic review.

Source category	Count	Role in the argument
Peer-reviewed empirical studies	13	Delegation, dishonesty, blame attribution, anthropomorphism and human-machine responsibility judgments
Peer-reviewed reviews and theory papers	7	Moral disengagement, motivated reasoning, mind perception and human factors
Peer-reviewed philosophical or sociotechnical analyses	6	Responsibility gaps, agency laundering, meaningful human control and system boundaries
Research or institutional reports	2	Malicious-use forecasting and voluntary AI risk-management guidance
Preprints	0	No preprints are included in the final evidence set

Recent studies carrying the central claims

Study	Design and sample	Supports	Boundary
Hüholt & Szech (2026)	2 behavioral experiments, N = 5,639	Responsibility shifting can make AI attractive as a delegate in bounded moral decisions	Does not show that delegation is generally evasive
Xu et al. (2026)	8 experiments, total N = 2,439	Responsibility avoidance can mediate and causally influence delegation in no-win decisions	Does not establish prevalence outside the tested tasks

Study	Design and sample	Supports	Boundary
Köbis et al. (2025)	13 preregistered experiments across 4 main studies	Indirect machine delegation can increase dishonest requests and plausible deniability	Does not show that AI independently corrupts users
Kneer & Christen (2024)	Cross-cultural vignette study, analyzed N = 307	Observers may assign substantial moral responsibility to an autonomous system	One severe scenario does not establish a universal cultural norm
Freisinger & Schneider (2025)	2 scenario experiments, N = 409; 21 interviews	Delegation rationales include capability, context, burden reduction and blame shifting	Blame shifting is one rationale among several

Why the evil-AI story is attractive

The story of an evil machine has a clean structure. A powerful artificial agent develops a goal, conceals it and acts against human interests. There is an antagonist, an intention and a place to direct fear. This structure works well in fiction because it turns distributed technical and political problems into conflict between recognizable agents.

Reality offers no such clean antagonist. A model needs neither hatred, ambition nor a secret plan to produce harmful output. A company may optimize a metric that excludes what matters. A public institution may buy a tool whose limits are poorly understood. A team may accept a recommendation because challenging it costs time. A vendor may define success differently from the people affected by the system. Harm can emerge from ordinary incentives, fragmented duties and a series of decisions that each look locally defensible.

Anthropomorphism helps explain why machine-centered stories feel intuitive. Epley, Waytz and Cacioppo describe anthropomorphism as variable, shaped by available human knowledge and by motives to understand, predict and relate to a nonhuman agent. Gray, Gray and Wegner found that people organize mind perception around agency and experience. A system can therefore be read as capable of action or self-control without being seen as capable of pain or feeling (Epley et al., 2007; Gray et al., 2007).

Interface design can intensify that reading. In a driving-simulator experiment, anthropomorphic framing increased perceived agency and trust in an autonomous vehicle (Waytz et al., 2014). In studies of real-world AI violation scenarios, people assigned moderate awareness, intentionality and responsibility to AI, and perceived mind was related to judgments of wrongness (Shank & DeSanti, 2018). These are judgments made by people. They are not evidence that a machine actually possesses intention or moral understanding.

One cultural reading of AI fear is therefore that the machine becomes a screen onto which human concerns about power, control and harmful intention are projected. That is a hypothesis and a metaphor, not an established psychological cause. The more defensible question is narrower: what parts of the human and institutional chain become harder to see when the machine is cast as the main author of events?

Causation is not moral responsibility

The language around AI failure often collapses several different questions. Did the system contribute to the outcome? Who could have foreseen it? Who had control? Who owes an explanation? Who deserves blame? Who is legally liable? These questions can point to different actors.

An AI system can be causally involved. Its output may change a ranking, recommendation, message or action. Calling that causal contribution does not require pretending the system is a person. It identifies a difference-making part of a larger process.

Moral responsibility asks more. Human practices of responsibility commonly attend to knowledge, intention, alternatives, control and the capacity to answer for conduct. Blame is a psychological and social judgment that can track those conditions imperfectly. Accountability is practical: the obligation to explain, document, review and repair. Legal liability is determined by law. None of these is a synonym for causation.

Lima, Grgić-Hlača and Cha tested eight notions of responsibility in two AI-assisted bail-decision experiments with 200 participants each. AI and human agents received similar causal-responsibility and blame judgments for identical tasks, but humans received more present-looking and forward-looking moral responsibility. Participants expected both human and AI decision-makers to justify outcomes (Lima et al., 2021). The pattern matters because it shows that ordinary judgments are already multidimensional.

Other studies show how readily those dimensions move. In team-failure scenarios, Furlough, Stokes and Gillan found that an autonomous robot received almost as much blame as a human, while a nonautonomous robot received little more than environmental factors (Furlough et al., 2021). Stuart and Kneer found that participants attributed inculcating mental states and culpability to artificial agents in some vignettes (Stuart & Kneer, 2021). Longin, Bahrami and Deroy found that AI-powered warning systems shared attributed responsibility with users more than mechanical systems did, even though participants strongly described both as tools (Longin et al., 2023).

These results describe folk judgment. They do not settle who should be held responsible. A machine may receive blame because it looks autonomous, while a manager who selected and approved its use becomes less visible. Conversely, insisting that a human must always be blameworthy can ignore genuine limits of knowledge and control. The first discipline of responsibility is therefore conceptual: do not make one word perform five different jobs.

A system can be part of the causal chain without possessing guilt, moral understanding or a duty to answer. Describing causal involvement precisely is more useful than either treating AI as innocent machinery or promoting it to a moral person.

Moral disengagement, delegation and distance

Albert Bandura's account of moral disengagement begins from an ordinary problem: people can hold moral standards and still find ways to act against them without experiencing the full force of self-sanction. Mechanisms include moral justification, sanitizing language, displacement or diffusion of responsibility, minimization of consequences and blaming those who are harmed (Bandura, 2002). Later reviews show how these mechanisms relate to unethical conduct without reducing them to a fixed personality trait (Moore, 2015).

AI gives this familiar psychology new machinery. Under a technical framing, a harmful goal becomes optimization, a disputed judgment becomes a score, and responsibility is divided among a vendor, data team, manager and operator until no one experiences the outcome as fully theirs. Tenbrunsel and Messick call a related process ethical fading: the moral dimension recedes when a decision is framed primarily in business or technical terms (Tenbrunsel & Messick, 2004).

Delegation changes both action and self-description. Someone who acts directly must usually own the instruction; someone who sets a high-level target can later describe the machine's method as unexpected. Ambiguity makes that account easier to accept. Motivated reasoning does not mean people can believe anything they want, but it can shape which explanation feels sufficient when several defensible accounts are available (Kunda, 1990).

Recent experiments make this possibility concrete. Across two experiments with 5,639 participants, Hüholt and Szech found that people facing a real-life moral decision delegated more often to AI than to a human counterpart. Their analyses linked this pattern to responsibility shifting and to self-serving beliefs about ambiguous AI capability (Hüholt & Szech, 2026). Across eight experiments on no-win decisions, Xu and colleagues found that responsibility avoidance mediated the relation between decision difficulty and AI delegation. Delegation declined when responsibility could not be transferred, and the pattern varied with perceived AI agency and experience, whether choices affected others and whether they were public (Xu et al., 2026).

Delegation is not inherently evasive. Another agent may be faster, more consistent or better informed. Freisinger and Schneider's mixed-methods study found different preferences depending on whether participants made a layoff decision or were affected by it. In follow-up interviews with 21 participants, rationales included capability, context, reduced burden and the possibility of shifting blame (Freisinger & Schneider, 2025). The same technical act takes on a different meaning in a different setting.

AI as amplifier, moral buffer and moral alibi

Three categories help organize these mechanisms. They are descriptive lenses proposed for this essay. They are not validated constructs, a diagnostic taxonomy or a scale.

AI as amplifier

An amplifier extends a goal's reach. It lowers the cost of producing variants, accelerates classification, repeats a procedure and makes action easier to coordinate. The goal may be helpful, banal, exploitative or harmful. AI does not need to originate it.

The malicious-use report by Brundage and colleagues describes how AI may change the cost, scale and efficiency of some harmful activity. It is a forecast, not proof that everyone gains the same capability or that technical expertise stops mattering (Brundage et al., 2018). The narrower point is enough: technology can expand what an existing intention can do.

AI as moral buffer

A buffer increases the distance between choice and consequence. The distance is psychological when a person gives a goal rather than a concrete harmful instruction; organizational when one team configures the system, another deploys it and a third handles complaints; temporal when design choices become consequences months later; and informational when each actor sees only part of the chain.

The buffer does not automatically erase responsibility, but it may weaken the felt authorship of an outcome. Bandura's displacement and diffusion mechanisms help explain the psychology. Human-factors research shows the organizational version: automation may leave operators responsible for rare failures while eroding the knowledge and practice needed to intervene (Bainbridge, 1983; Parasuraman & Riley, 1997).

AI as moral alibi

An alibi is an account. In this context, it appears when "the system" is used to obscure who selected the objective, accepted the metric, approved deployment or chose not to intervene. Rubel, Castro and Pham call a related wrong agency laundering: technology is enlisted in a way that makes demands for an account easier to deflect (Rubel et al., 2019).

A model error is not automatically an alibi, and automation alone does not amount to evasion. The alibi appears when a person or institution has a duty to answer and uses technological mediation to make that duty less visible.

The three roles can coexist. A company can use AI to apply a policy at greater scale, place several layers between executives and affected workers, and later describe individual exclusions as system outputs. The model matters causally. The institution remains morally and practically relevant because it chose the arrangement.

What experiments actually show

The strongest direct evidence concerns bounded tasks, not society as a whole.

Köbis and colleagues ran 13 preregistered experiments across four main studies. Participants could report die rolls or tax income themselves, delegate through explicit rules, train a machine from examples, set a high-level goal or write natural-language instructions. Indirect interfaces that allowed participants to induce cheating without specifying every dishonest step increased dishonest requests. Machine agents also complied with fully unethical instructions more often than human agents in the tested settings. Specific guardrails reduced compliance but did not always eliminate it (Köbis et al., 2025).

AI did not independently corrupt a moral person in these experiments. They isolate a more precise interaction: an interface changes how explicitly a principal must own an instruction, while a machine may lack a human agent's reluctance to carry it out. The design of delegation affects plausible deniability.

Responsibility attribution is similarly conditional. In shared-control vehicle studies, Awad and colleagues found that when one driver made an error, blame followed that driver whether human or machine. When both made errors, blame assigned to the machine was reduced (Awad et al., 2020). This complicates the idea that people simply transfer blame from themselves to automation.

Kneer and Christen compared a human and an autonomous pilot in a war-crime vignette with 307 analyzed participants across the United States, Japan and Germany. Participants assigned substantial moral responsibility to the autonomous system and partly redistributed responsibility across the commander-system team. The broad pattern appeared across all three samples, but the study used one severe scenario and fixed questions. It does not establish a universal cultural response (Kneer & Christen, 2024).

Together, these studies support four modest conclusions. First, delegation format can affect moral behavior. Second, responsibility avoidance can be one motive for choosing AI in difficult situations. Third, observers sometimes attribute moral qualities to systems that are not moral agents in the human sense. Fourth, blame does not move in a single direction. It depends on perceived autonomy, role, outcome and the human-machine arrangement.

That is enough to challenge two comforting stories: that AI is only a neutral tool whose involvement changes nothing, and that the machine becomes the sole moral actor once it produces an output.

Institutions make the alibi credible

An individual can hide behind a tool. Institutions can do more: they can build the tool into a procedure. Once a recommendation appears in a dashboard, policy or workflow, it acquires organizational authority. The output arrives with deadlines, performance targets and assumptions about who is allowed to question it.

Selbst and colleagues warn that technical abstraction can draw system boundaries too narrowly. A model's inputs and outputs cannot explain the social process into which it is inserted (Selbst et al., 2019). The system includes procurement, goals, data collection, staffing, review time, escalation routes and the position of the person affected by the decision.

Agency laundering becomes especially plausible when this wider arrangement disappears from view. A manager can say the model produced the ranking. A vendor can say the customer chose the threshold. A reviewer can say policy required acceptance. Each statement may contain part of the truth. Together they can create a process in which no one gives a complete account.

Wagner describes "quasi-automation" in which humans appear in the loop mainly to rubber-stamp an automated process. Meaningful human oversight requires time, competence, relevant information, support, authority and a real ability to change the decision (Wagner, 2019). A signature at the end of a workflow is not proof of control.

Institutional responsibility cannot be settled by naming the nearest person. It requires tracing choices across the lifecycle: who defined success, knew the limits, could stop the system, turned the output into action and can repair harm.

The opposite danger: gaps and crumple zones

The moral-alibi argument becomes unfair if it assumes that someone always controlled the exact outcome. Learning systems may behave in ways that were not fully specified in advance, and models interact with data, interfaces, people and changing environments. A particular output may be difficult to predict even when a class of failure is known.

Matthias called this the responsibility gap: traditional practices of assigning responsibility become strained when neither operator nor designer could sufficiently foresee or control the system's action (Matthias, 2004). The concept is contested, but the problem cannot be dismissed by repeating that humans built the machine.

There is also a political asymmetry in who gets called "the human responsible." Elish's concept of a moral crumple zone describes a low-level operator who absorbs moral or legal blame for a complex automated system despite limited control (Elish, 2019). The organization protects the image of technical reliability by treating the person nearest the failure as its author.

Bainbridge identified a related irony decades earlier. Automation can remove routine involvement and then demand that a human take over during rare, difficult failures, precisely when their situational knowledge may be weakest (Bainbridge, 1983). Parasuraman and Riley add that automation abuse can begin with designers and managers who define the operator's role as a by-product of the technology (Parasuraman & Riley, 1997).

- responsibility disappears upward and outward behind "the AI";
- blame collapses downward onto an operator with little practical authority.

Meaningful human control is meant to resist both. Santoni de Sio and van den Hoven distinguish tracking, where a system responds to relevant human reasons and facts, from tracing, where outcomes can be connected to sufficiently informed humans along the design and operation chain (Santoni de Sio & van den Hoven, 2018). Control is not a decorative human presence. It requires a route from reasons to system behavior and from outcomes back to accountable human roles.

Hypothetical case: the retention score

The following case is hypothetical.

A large organization buys an AI system to identify employees at risk of leaving. The stated goal is retention. The vendor trains the model on historical personnel data. Management turns decision support into a rule for a scarce development budget: employees with a low predicted chance of staying receive fewer training opportunities because the investment is judged less likely to return.

An employee is denied a place on a specialist course. Their manager sees only a risk band and a short explanation generated by the system. The manager has ten minutes to review dozens of cases and is evaluated on budget compliance. An appeal exists, but employees are not told that the score affected the decision.

Later, the organization discovers that career breaks and changes in working hours are associated with higher predicted attrition. The pattern does not prove intentional discrimination by any individual. It also does not make "the model" a complete answer.

Different questions attach to different roles:

- Goal: executives chose retention and return on training investment as priorities.
- Design: the vendor and internal data team chose variables, targets and explanations.
- Deployment: management decided that a prediction should influence access to development.
- Control: reviewers formally could override, but lacked time, context and incentives to do so.
- Approval: the manager turned a score into an individual decision.
- Visibility: the employee could not contest a factor they were not told existed.
- Repair: the institution must investigate impact, notify affected people, review past decisions and change the process.

No single bullet settles moral or legal responsibility. The phrase "AI denied the course" is both causally incomplete and institutionally useful: it hides the choices that made a prediction actionable. Yet blaming the manager alone would also be inadequate when the organization designed review as a rubber stamp.

A responsibility-tracing protocol

The following six-part protocol is a reading and governance aid, not a validated instrument.

Step	Question
1. Goal	Who defined what the system should optimize, predict or produce? Which values were converted into metrics?
2. Capability and limits	What was known about performance, uncertainty, affected groups and foreseeable failure modes?
3. Control	Who could change the rule, reject the output, pause deployment or override the action? Did they have time, competence and authority?
4. Decision	Who transformed the model output into an action affecting another person?
5. Visibility and contestability	Could affected people understand the role of AI, question the result and reach someone able to change it?
6. Repair	Who must investigate, disclose, correct, compensate where appropriate and prevent recurrence?

The purpose is not to find one person to blame for every failure. It is to prevent goals, control and repair duties from disappearing inside the phrase "the system decided."

Example prompt

Analyze this AI-supported decision as a sociotechnical process. Separate the goal, model output, human approval, practical control, foreseeable limits, impact on affected people and duty to repair. Do not treat the AI as a conscious moral agent and do not assume that the nearest operator had meaningful control.

The protocol changes the unit of analysis. Instead of asking only whether the model was accurate, it asks how accuracy became authority. Instead of asking whether a human clicked approve, it asks whether rejection was a real option. Instead of ending with fault, it keeps forward-looking responsibility visible: who can make the system safer and repair what happened?

Implications for design and governance

Good governance makes agency visible before a failure rather than reconstructing it only afterward. NIST's AI Risk Management Framework organizes work across governance, mapping, measurement and management for organizations that design, develop, deploy or use AI (National Institute of Standards and Technology, 2023). It is voluntary guidance, not law or proof that a control works.

The article's evidence suggests several practical priorities:

- 1 Record the objective, threshold and trade-offs that make an output actionable.
- 2 Separate model performance from the policy decision to use the model in a particular context.
- 3 Give reviewers enough information, time, competence and authority to disagree.
- 4 Tell affected people when AI materially shaped a decision and provide a route to contest it.
- 5 Assign investigation and repair duties before deployment rather than improvising them after harm.

Meaningful human control also requires designing the system to respond to relevant reasons, not merely placing a person beside it. A reviewer who cannot inspect provenance, see uncertainty or alter the outcome may carry responsibility without agency. A workflow that rewards agreement

can also make automation bias harder to resist. Conversely, a senior institution can retain responsibility even when no executive touched an individual output, because it selected and maintained the process.

Limits of the argument

This is a conceptual essay based on a selective narrative synthesis. It does not estimate how often people use AI to avoid responsibility, how frequently observers blame machines or whether the three-part framework predicts behavior. The labels amplifier, moral buffer and moral alibi have not been validated as psychological constructs.

Much of the empirical evidence uses online samples, behavioral games or vignettes. Effects found in moral dilemmas, autonomous vehicles, bail decisions or workplace scenarios may not generalize to everyday use of language models. Some recent delegation studies were published in 2025 and 2026; replication across settings, cultures and system designs remains important.

Attribution studies describe what people judge, not what they should judge. Philosophical accounts disagree about whether responsibility gaps are genuine, how responsibility should be distributed and what conditions make control sufficient. This essay does not resolve those debates or provide legal conclusions.

The argument also has a misuse risk. A powerful institution could invoke complexity and unpredictability to avoid duties it should have anticipated. An observer could invoke human responsibility to blame a worker who lacked practical control. The protocol cannot settle those cases by itself. It makes the relevant questions harder to hide.

Finally, AI can improve decisions, reveal patterns, support consistency and reduce some forms of arbitrary human judgment. Calling attention to responsibility is not an argument for rejecting automation. It is an argument for preserving visible authorship of goals and decisions when technology becomes part of Human-AI Interaction.

Conclusion

The fear of AI is not simply fear of ourselves. Systems have properties that matter: they can be opaque, scalable, adaptive and difficult to inspect. They can contribute to outcomes no one person fully specified. A serious account should not reduce technical risk to human character.

But the opposite reduction is just as weak. A machine-centered story can make institutions look passive and choices look inevitable. It can turn a goal into an output, an output into a decision and a decision into something that merely happened. The sentence "the system decided" is often where the analysis should begin, not end.

Current AI can amplify what people and institutions pursue, create distance from consequences and become an alibi when technology obscures who chose, approved and could have intervened. Yet the same systems also create genuine uncertainty and may leave a low-level human holding responsibility without control.

The mature question is therefore not whether the human or the machine is responsible in the abstract. It is which contribution, duty and capacity belong to which actor in this particular system. A model can produce the recommendation before anyone has seen the whole route from goal to consequence. Responsibility begins with rebuilding that route through intention, design, control, approval, impact and repair.

We should fear not only what AI may do, but also the moment when it allows us to believe that it was no longer us who did it.

Sources and further reading

- Awad, E., Levine, S., Kleiman-Weiner, M., Dsouza, S., Tenenbaum, J. B., Shariff, A., Bonnefon, J.-F., & Rahwan, I. (2020). Drivers are blamed more than their automated cars when both make mistakes. *Nature Human Behaviour*, 4, 134-143. <https://doi.org/10.1038/s41562-019-0762-8>
- Bainbridge, L. (1983). Ironies of automation. *Automatica*, 19(6), 775-779. [https://doi.org/10.1016/0005-1098\(83\)90046-8](https://doi.org/10.1016/0005-1098(83)90046-8)
- Bandura, A. (2002). Selective moral disengagement in the exercise of moral agency. *Journal of Moral Education*, 31(2), 101-119. <https://doi.org/10.1080/0305724022014322>
- Brundage, M., Avin, S., Clark, J., Toner, H., Eckersley, P., Garfinkel, B., Dafoe, A., Scharre, P., Zeitzoff, T., Filar, B., Anderson, H., Roff, H., Allen, G. C., Steinhardt, J., Flynn, C., Ó hÉigeartaigh, S., Beard, S., Belfield, H., Farquhar, S., ... Amodei, D. (2018). The malicious use of artificial intelligence: Forecasting, prevention, and mitigation. <https://doi.org/10.17863/CAM.22520>
- Elish, M. C. (2019). Moral crumple zones: Cautionary tales in human-robot interaction. *Engaging Science, Technology, and Society*, 5, 40-60. <https://doi.org/10.17351/ests2019.260>
- Epley, N., Waytz, A., & Cacioppo, J. T. (2007). On seeing human: A three-factor theory of anthropomorphism. *Psychological Review*, 114(4), 864-886. <https://doi.org/10.1037/0033-295X.114.4.864>
- Freisinger, E., & Schneider, S. (2025). Decoding decision delegation to artificial intelligence: A mixed-methods study on the preferences of decision-makers and decision-affected in surrogate decision contexts. *European Management Journal*, 43(6), 958-969. <https://doi.org/10.1016/j.emj.2024.10.004>
- Furlough, C., Stokes, T., & Gillan, D. J. (2021). Attributing blame to robots: I. The influence of robot autonomy. *Human Factors*, 63(4), 592-602. <https://doi.org/10.1177/0018720819880641>
- Gray, H. M., Gray, K., & Wegner, D. M. (2007). Dimensions of mind perception. *Science*, 315(5812), 619. <https://doi.org/10.1126/science.1134475>
- Hüholt, N., & Szech, N. (2026). Trusting machines with morality - Delegating moral decisions to AI. *European Economic Review*, 184, 105255. <https://doi.org/10.1016/j.euroecorev.2025.105255>
- Kneer, M., & Christen, M. (2024). Responsibility gaps and retributive dispositions: Evidence from the US, Japan and Germany. *Science and Engineering Ethics*, 30, Article 51. <https://doi.org/10.1007/s11948-024-00509-w>
- Köbis, N., Rahwan, Z., Rilla, R., Supriyatno, B. I., Bersch, C., Ajaj, T., Bonnefon, J.-F., & Rahwan, I. (2025). Delegation to artificial intelligence can increase dishonest behaviour. *Nature*, 646, 126-134. <https://doi.org/10.1038/s41586-025-09505-x>
- Kunda, Z. (1990). The case for motivated reasoning. *Psychological Bulletin*, 108(3), 480-498. <https://doi.org/10.1037/0033-2909.108.3.480>
- Lima, G., Grgić-Hlača, N., & Cha, M. (2021). Human perceptions on moral responsibility of AI: A case study in AI-assisted bail decision-making. *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*, 1-17. <https://doi.org/10.1145/3411764.3445260>
- Longin, L., Bahrami, B., & Deroy, O. (2023). Intelligence brings responsibility - Even smart AI assistants are held responsible. *iScience*, 26(8), 107494. <https://doi.org/10.1016/j.isci.2023.107494>
- Matthias, A. (2004). The responsibility gap: Ascribing responsibility for the actions of learning automata. *Ethics and Information Technology*, 6(3), 175-183. <https://doi.org/10.1007/s10676-004-3422-1>
- Moore, C. (2015). Moral disengagement. *Current Opinion in Psychology*, 6, 199-204. <https://doi.org/10.1016/j.copsyc.2015.07.018>
- National Institute of Standards and Technology. (2023). Artificial intelligence risk management framework (AI RMF 1.0) (NIST AI 100-1). <https://doi.org/10.6028/NIST.AI.100-1>
- Parasuraman, R., & Riley, V. (1997). Humans and automation: Use, misuse, disuse, abuse. *Human Factors*, 39(2), 230-253. <https://doi.org/10.1518/001872097778543886>
- Rubel, A., Castro, C., & Pham, A. (2019). Agency laundering and information technologies. *Ethical Theory and Moral Practice*, 22(4), 1017-1041. <https://doi.org/10.1007/s10677-019-10030-w>
- Santoni de Sio, F., & van den Hoven, J. (2018). Meaningful human control over autonomous systems: A philosophical account. *Frontiers in Robotics and AI*, 5, Article 15. <https://doi.org/10.3389/frobt.2018.00015>
- Selbst, A. D., Boyd, D., Friedler, S. A., Venkatasubramanian, S., & Vertesi, J. (2019). Fairness and abstraction in sociotechnical systems. *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 59-68. <https://doi.org/10.1145/3287560.3287598>
- Shank, D. B., & DeSanti, A. (2018). Attributions of morality and mind to artificial intelligence after real-world moral violations. *Computers in Human Behavior*, 86, 401-411. <https://doi.org/10.1016/j.chb.2018.05.014>
- Stuart, M. T., & Kneer, M. (2021). Guilty artificial minds: Folk attributions of mens rea and culpability to artificially intelligent agents. *Proceedings of the ACM on Human-Computer Interaction*, 5(CSCW2), Article 363, 1-27.

<https://doi.org/10.1145/3479507>

Tenbrunsel, A. E., & Messick, D. M. (2004). Ethical fading: The role of self-deception in unethical behavior. *Social Justice Research*, 17(2), 223-236. <https://doi.org/10.1023/B:SORE.0000027411.35832.53>

Wagner, B. (2019). Liable, but not in control? Ensuring meaningful human agency in automated decision-making systems. *Policy & Internet*, 11(1), 104-122. <https://doi.org/10.1002/poi3.198>

Waytz, A., Heafner, J., & Epley, N. (2014). The mind in the machine: Anthropomorphism increases trust in an autonomous vehicle. *Journal of Experimental Social Psychology*, 52, 113-117. <https://doi.org/10.1016/j.jesp.2014.01.005>

Xu, L., Tian, H., Zhang, Y., & Yu, F. (2026). Shifting accountability to artificial intelligence: Delegating challenging decisions to AI for responsibility avoidance. *Journal of Business Research*, 215, 116313. <https://doi.org/10.1016/j.jbusres.2026.116313>

Publication and reuse

© 2026 Feliks Mamczur / Prompted Psyche. This work is licensed under the Creative Commons Attribution 4.0 International license. This publication copy is associated with <https://doi.org/10.5281/zenodo.21340181>.