

Beyond AI Share: A Preregistered Survey and Vignette Study of Perceived Control, Authorship, and Authenticity in AI-Assisted Creative Practice

Feliks Mamczur

Independent researcher

Warsaw, Poland

mamczur@gmail.com

Preprint

Version 1.0

Not peer reviewed

Publication date: 2026-07-30

DOI: <https://doi.org/10.5281/zenodo.21705721>

© 2026 Feliks Mamczur

Licensed under Creative Commons Attribution 4.0 International

Suggested citation:

Mamczur, F. (2026). Beyond AI Share: A Preregistered Survey and Vignette Study of Perceived Control, Authorship, and Authenticity in AI-Assisted Creative Practice [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21705721>

Abstract

AI involvement in creative work can be described by quantity, but quantity alone does not reveal how decisions are organized or whether human input remains consequential. This preregistered cross-sectional online survey with two within-person vignettes examined perceived authorship, authenticity, and control in AI-assisted creative practice. The full sample comprised 429 adults, including 164 creators who used AI. Among these creators, declared AI share was not negatively associated with perceived authorship ($r = .141$, $p = .071$), contrary to H1. Perceived control was positively associated with authorship ($r = .234$, $p = .003$), supporting H2. Control did not moderate the association between AI share and authenticity (interaction $b = .043$, $p = .559$), so H3 was not supported. In the full sample, expressive orientation toward art was associated with creative identity threat ($r = .341$, $p < .001$), cautiously supporting H4. The strongest result was the vignette contrast: a human-directed process involving idea formation, selection, and substantial revision was evaluated more favorably than acceptance of a near-final AI output (mean difference = .836, $t = 13.628$, $p < .001$, $d_z = .659$). Several short indicators had low reliability, limiting strong inference. Meaningful human control is therefore used only as an interpretive lens, not as a validated model. The findings suggest that digital creative practice and tool design should attend to process structure and consequential human control rather than relying on AI-share estimates alone.

Keywords: human-AI co-creativity; creative practice; authorship; authenticity; perceived control; generative AI

1 Introduction

Generative artificial intelligence is becoming part of ordinary creative work. Writers use language models to explore phrasing or narrative alternatives; designers generate visual variants; musicians test arrangements and timbres; photographers and filmmakers use AI in planning, editing, and post-production. These practices differ greatly, yet they are often compressed into the same label: “AI-assisted.” The label identifies the presence of a technology but says little about how a work was made, which decisions remained human, or how the creator relates the final artifact to their own intention and judgment.

Debates about AI-assisted creativity frequently rely on binary distinctions. A work is described as human-made or AI-generated, authentic or inauthentic, original or derivative. Binary categories can be useful for disclosure, but they do not represent a workflow in which human and machine contributions are distributed across ideation, generation, evaluation, editing, and final approval. Nor does a percentage estimate of AI involvement identify which contribution determined the concept, composition, narrative, or final selection. A creator may use extensive generated material while directing the concept, rejecting most outputs, combining selected elements, and substantially revising the result. Another may enter a brief, accept the first polished output, and make only cosmetic changes. Both practices can receive the same AI-assisted label and a similar declared AI share, although their structures of agency differ.

The distinction matters because the study concerns psychological authorship and authenticity, not only technical contribution. Perceived authorship refers to the extent to which a creator experiences or is judged to

experience the work as their own. Authenticity concerns whether the work is seen as preserving human intention, experience, or personal style. Perceived control concerns whether the person can direct consequential decisions rather than merely confirm an output. These constructs may be related, but none is reducible to the quantity of generated material.

Meaningful human control offers one interpretive lens for describing this process. In autonomous-systems scholarship, the concept has been developed through tracking and tracing conditions and through actionable properties for system development (Santoni de Sio and van den Hoven 2018; Cavalcante Siebert et al. 2023). Here, direction, selection, revision, and responsibility are used as a domain-specific adaptation for creative practice. They organize questions about whether human involvement remains consequential to the development and final form of a work. The present study did not measure these four elements as separate latent variables and does not validate a universal framework or scale of meaningful human control.

Research across creative domains points toward distinctions among roles, stages, and allocations of initiative, including a shift from direct production toward curation and meta-creation while human judgment remains important in later phases (Tsao et al. 2026). Attribution research asks which contributions merit credit (He et al. 2025), while studies of AI-assisted writing identify gaps among formal authorship, perceived ownership, and authenticity (Draxler et al. 2024; Hwang et al. 2025). Artist-focused work places transparency, fairness, professional identity, artistic control, and workplace conditions within the same debate (Lovato et al. 2024; Weaver 2026; Jiang et al. 2026). In education, agency and ethical awareness are likewise tied to creator identity (Lee et al. 2026). Together, these strands suggest that output quality, AI presence, and AI-share estimates cannot be the only criteria for understanding creative experience.

The study reported here combines a preregistered survey with two within-person vignettes. The self-report component examines declared AI involvement, perceived control, authorship, authenticity, expressive orientation, and creative identity threat. The vignette component contrasts a workflow in which a creator begins with an idea, requests variants, selects one, revises it substantially, and determines the final composition with a workflow in which a creator accepts a near-final AI output after only minor changes. The correlational measures test associations in reported experience; the bundled vignette contrast tests whether respondents distinguish two process structures. Neither component establishes a causal standard for creative practice.

The central argument is that the declared quantity of AI involvement is insufficient to explain perceived authorship and authenticity. Evaluations also depend on how the creative process is structured and on the degree of consequential human control retained within it. This claim is tested cautiously: the cross-sectional survey uses several short, partly unreliable indicators, and the vignettes bundle multiple process features. The contribution is therefore a bounded comparison between quantity-based and process-sensitive explanations, not a definitive theory of authorship or a validated model of control.

2 AI involvement and the structure of creative contribution

2.1 The AI-assisted label is informationally thin

Human-AI co-creativity encompasses systems that execute tightly specified instructions, suggest alternatives, engage in iterative exchange, or generate near-final artifacts. The person may originate the concept, curate machine outputs, revise components, combine materials, or simply authorize delivery. A systematic account must therefore distinguish stages, roles, and initiative rather than treat AI as a uniform collaborator (Tsao et al. 2026; Haidemariam 2026).

This variation limits what can be inferred from disclosure alone. Saying that AI was used can be important for transparency, but it does not reveal the significance of the machine contribution. A short generated suggestion may redirect an entire composition. Conversely, a large quantity of generated material may remain subordinate to a clearly specified human concept and extensive transformation. A percentage of “AI contribution” can communicate perceived quantity, but it cannot fully represent the importance of decisions made by the system or by the person.

2.2 Contribution is qualitative as well as quantitative

Attribution depends on more than participation. Research on credit in human-AI co-creation suggests that contributions are judged according to their character and importance, including whether they frame the goal, introduce a distinctive idea, or determine the final selection (He et al. 2025). This insight is familiar from human collaboration, where assistance, delegation, and specialized tools do not automatically erase authorship. Credit often follows the organizing contribution that gives other inputs their role.

The distinction between authorship and ownership illustrates the problem. Draxler et al. (2024) describe an AI ghostwriter effect in which people may declare themselves authors without fully perceiving ownership of AI-generated text. Public attribution, procedural responsibility, and subjective attachment can diverge. A person may sign and deliver a work while feeling that its wording or form does not belong to them. A process-sensitive account does not assume that any single dimension settles the issue; it asks how intention, decision, and outcome remain connected.

2.3 Creative contribution is situated in tools, organizations, and norms

Creative AI is used within deadlines, contracts, educational settings, platform defaults, and professional hierarchies. A tool may technically permit revision while a workflow rewards immediate acceptance. An employee may be formally responsible for an output but lack authority to refuse AI use. A student may be instructed to use a model without learning how to evaluate its suggestions. Meaningful control therefore has both interface-level and organizational conditions.

Artist-centered studies highlight transparency, ownership, and fairness as practical concerns (Lovato et al. 2024). Practice-based research in music shows that generative AI can expand creative materials while renegotiating artistic control and perceived collaboration (Weaver 2026). A cross-domain review identifies

hybrid methods that often position generative AI in early ideation while retaining human judgment in convergent phases (Tsao et al. 2026). Research with professional visual artists similarly places adoption within workplace negotiation and changing expectations about labor (Jiang et al. 2026). These findings caution against equating use with endorsement or non-use with resistance. A creator can regard AI as useful while worrying about professional identity, compensation, or the erosion of acquired skills.

The creativity-support literature offers a complementary perspective: tools should be assessed not only by efficiency but also by their support for exploration, expression, and work worth producing (Cherry and Latulipe 2014; Cox et al. 2025). For generative AI, speed and variation can coexist with premature closure or reduced agency. Evaluation must therefore ask both what capacities a tool expands and what decisions it displaces.

3 Perceived control, authorship, and authenticity

Meaningful human control is used here as a process-level interpretive lens, not as the name of a validated scale. It concerns whether human involvement is consequential to the development and final form of a work. A superficial confirmation step does not necessarily provide meaningful control, just as extensive manual labor does not guarantee authorship or authenticity. Control is meaningful in this bounded sense when the person can shape the workflow, intervene in relevant decisions, and remain answerable for the outcome.

Meaningful human control was originally formulated through tracking and tracing conditions (Santoni de Sio and van den Hoven 2018) and later translated into actionable properties for AI system development (Cavalcante Siebert et al. 2023). Tracking concerns whether system behavior responds to relevant human reasons and environmental facts; tracing concerns whether outcomes can be connected to at least one responsible human in the design or use chain. Control-by-design work further emphasizes multilevel feedback and continuing human learning rather than nominal human presence alone (Te’eni et al. 2026). Direction, selection, revision, and responsibility are a domain-specific adaptation for creative practice, not a new universal definition or a validated scale.

The lens has four analytically distinct but connected components: direction, selection, revision, and responsibility. They are not claimed to form a complete taxonomy. They organize the aspects of human-directed practice represented in the present measures and vignettes and provide a basis for interpreting the results.

3.1 Direction

Direction concerns the formation of purpose. The creator establishes what the work is trying to do, which constraints matter, what audience or context is relevant, and how success will be evaluated. A prompt can express direction, but prompt entry alone is not sufficient evidence of it. A vague request followed by uncritical acceptance differs from an iterative process guided by a stable intention.

3.2 Selection

Selection concerns the ability to compare, accept, reject, and combine alternatives. Generative systems often shift effort from producing a single candidate to navigating many candidates. This can expand exploration, but it can also reduce the user's role to browsing outputs optimized by the system. Selection is meaningful when alternatives are evaluated against human criteria and when rejection is a genuine option.

3.3 Revision

Revision concerns whether the person can materially transform generated content. Editability matters because it turns an output into an object of continued creative work. Revision may involve rewriting, recomposing, masking, rearranging, combining sources, changing parameters, or returning to an earlier stage. It is meaningful when these actions can alter features central to the work rather than merely correct superficial defects.

3.4 Responsibility

Responsibility concerns the person's capacity to endorse, explain, and remain accountable for the final work. It does not require complete technical knowledge of a model, just as responsibility for conventional tools does not require knowledge of every underlying mechanism. It does require enough understanding of the workflow to make informed decisions about suitability, risks, disclosure, and delivery.

3.5 Control without total predictability

Meaningful control is not equivalent to deterministic control. Creative systems are often valuable because they generate novelty and surprise. Nor does the framework imply that the human must perform every operation manually. Assistance can be extensive while remaining organized by human purposes. The key question is whether the person retains effective opportunities to guide, evaluate, and alter the path from intention to output.

Perceived control is related to the sense of agency, or the experience of initiating and directing action (Tapal et al. 2017). The present study measures it through a short indicator and process judgments through vignettes; it does not measure all four components as separate latent variables. The four-part account is therefore an interpretation to be refined, not a psychometrically established model.

3.6 Perceived authorship

Formal authorship, copyright, social attribution, and the subjective feeling that a work is one's own are related but not identical. The present study concerns perceived authorship: whether AI-assisted work is experienced or judged as belonging to the human creator. It does not determine legal authorship. This distinction is important because the same artifact can support different judgments about legal rights, professional credit, personal ownership, and responsibility.

Experimental vignette research indicates that the degree of assistance affects judgments of human authorship, creatorship, and responsibility more consistently than whether the assistance comes from a

person or AI (Formosa et al. 2025). Creative ownership accounts likewise emphasize the structure and significance of contributions (Polimetla et al. 2026). Authorship may be more plausible when a person directs, selects, materially revises, and accepts responsibility. Declared AI share remains an incomplete proxy because it does not distinguish execution from decision-making.

3.7 Authenticity

The idea of a “semi-aura” describes authenticity in AI-generated art as redistributed across human intention, technological mediation, and reception rather than simply present or absent (Salas Espasa and Camacho 2025). Research on co-writing with large language models similarly locates authenticity in negotiation between machine contribution and personal voice (Hwang et al. 2025). These accounts fit a process-sensitive approach: AI assistance does not automatically eliminate authenticity, but authenticity may depend on how the person’s intentions and revisions remain legible in the work.

The present study operationalizes authenticity as the perceived capacity of AI-assisted work to preserve human intention and personal style. A separate auxiliary item concerns a recognizable “typical AI aesthetic.” This distinction allows concern about standardized surface conventions to be examined without assuming that every AI-assisted process produces the same authenticity judgment.

4 Creative identity and attitudes toward AI

4.1 Creative identity threat

AI may threaten creative identity through concerns about the economic value of labor, the distinctiveness of learned competence, the future of a profession, or the meaning of being a creator. Domain-specific threat research shows that perceived threats from AI vary by application and function (Kieslich et al. 2021). In creative practice, threat is not reducible to dislike or non-use. A creator may find AI useful while fearing that institutions will devalue human work.

4.2 Expressive orientation toward art

An expressive orientation treats art primarily as an expression of human experience, intention, and sensibility. People who hold this orientation strongly may be especially sensitive to workflows that appear to substitute generated form for lived experience. This does not establish a correct definition of art. It identifies one value orientation that may explain differences in authenticity concerns and identity threat.

4.3 AI usefulness and AI as a tool

AI usefulness is conceptually separate from perceived control. A system can accelerate drafting, widen exploration, and expand available techniques while still encouraging unreflective acceptance. Separating usefulness from control avoids a simplistic division between enthusiasts and opponents. It also allows the coexistence of perceived benefit and identity threat to be examined without treating that combination as a validated ambivalence scale.

5 Current study and preregistered hypotheses

The study asked: How is the role of AI in the creative process associated with perceived authorship, perceived authenticity, perceived control, and creative identity threat?

The preregistered model proposed that AI involvement would relate to authorship and authenticity, with perceived control serving as a central condition. It also proposed that expressive orientation toward art would be associated with stronger concerns about AI and creative identity. The design combined self-report measures with two short vignettes. The self-report component captured respondents' general attitudes and, for users, experiences of AI-assisted work. The vignettes held the broad task constant while contrasting a human-directed process with acceptance of a near-final AI output.

Four hypotheses were preregistered as confirmatory:

H1. Among creators using AI, a greater declared AI share in the final creative output would be associated with lower perceived authorship.

H2. Among creators using AI, higher perceived control over the creative process would be associated with higher perceived authorship.

H3. Perceived control would moderate the association between AI final share and authenticity, such that the predicted negative association would be weaker at higher levels of control.

H4. Stronger expressive orientation toward art would be associated with stronger concern about authenticity and/or greater creative identity threat in the context of AI-assisted creation.

Two additional hypotheses were treated as exploratory. **H5** concerned the possible coexistence of high perceived AI usefulness and high creative identity threat, especially among more engaged creators. **H6** proposed that more frequent AI use would be associated with a stronger tendency to view AI as a tool or assistant rather than as an independent creator.

The vignette comparison was planned as a separate analysis. The human-controlled scenario described a creator with an original idea who used AI to produce background variants, selected one, edited it substantially, and determined the final composition. The final-generator scenario described a creator who entered a short brief, selected the first generated result, made only minor changes, and delivered it as the final work. Respondents evaluated authorship, authenticity, and control for each scenario.

The hypotheses and analytical plan remained unchanged from the preregistration. The process-sensitive framing guides interpretation but does not alter the confirmatory analyses.

6 Method

6.1 Design

The study used a cross-sectional online survey with a within-person vignette component. It examined associations among self-reported variables and differences in judgments between two scenarios. It was not a

causal experiment and was not designed as a full psychometric validation study. The questionnaire was administered in Polish using a frozen public version labeled `questionnaire_v1.0`.

6.2 Participants

Data for the main study were collected from 15 June to 5 July 2026. Eligibility required age 18 or older, informed consent, and completion of the questionnaire to a usable extent. Pilot responses were kept separate and were not included in the main dataset.

The final analytic sample contained 429 completed and consenting responses. Of these, 209 respondents (48.7%) were classified as creators or creative professionals and 220 (51.3%) as people who did not engage in creative work and primarily identified as art audiences. The creator category included professional and semi-professional creators, amateurs, students in artistic or creative fields, and people employed in creative work without primarily identifying as artists.

Two hundred and forty-seven respondents (57.6%) reported using AI at least occasionally, while 182 (42.4%) reported never using it. The key subgroup of creators using AI contained 164 respondents, exceeding the preregistered minimum of 150. A more active subgroup of 102 creators reported using AI several times per month or more often.

The sample included 237 women (55.2%), 180 men (42.0%), seven non-binary respondents (1.6%), two respondents using another description, one who preferred not to answer, and two missing responses. Age was measured categorically: 10.3% were aged 18-24, 14.5% were 25-34, 28.7% were 35-44, 21.0% were 45-54, and 24.0% were 55 or older; 1.4% preferred not to answer and one response was missing. Creative fields were heterogeneous, with the largest creator groups in photography, music or sound, film or video, graphic design and illustration, and writing or literature. Small cell sizes prevented confirmatory field-specific comparisons.

Recruitment used non-probability online outreach through social, professional, and creative networks. The sampling strategy was intended to reach both creators with AI experience and a comparison group, not to estimate population prevalence. The sample should therefore be understood as a convenience sample rather than representative of all creators or all adults. The planned recruitment targets were $N = 500$ overall, a minimum usable $N = 300$, and a minimum $N = 150$ in the key subgroup of creators using AI. Reliable counts of all survey views, starts, and incomplete submissions were not retained, so view, participation, and completion rates could not be calculated. The final $N = 429$ denotes completed, consenting, and retained records rather than a 100% completion rate.

Table 1. Sample composition

Characteristic	Category	n	%
Analytic sample	Total	429	100.00
Creative status	Creators and creative professionals	209	48.72
Creative status	Non-creators / primarily art audiences	220	51.28
AI use	Uses AI at least occasionally	247	57.58
AI use	Never uses AI	182	42.42
Key subgroup	Creator AI users	164	38.23
Key subgroup	Active creator AI users	102	23.78
Gender	Women	237	55.24
Gender	Men	180	41.96
Gender	Non-binary	7	1.63
Gender	Other, prefer not to answer, or missing	5	1.17
Age	18-24	44	10.26
Age	25-34	62	14.45
Age	35-44	123	28.67
Age	45-54	90	20.98
Age	55 or older	103	24.01
Age	Prefer not to answer or missing	7	1.63

6.3 Measures

Unless otherwise stated, attitudinal items used a five-point agreement scale from 1 (strongly disagree) to 5 (strongly agree). Indicator scores were calculated as item means after reverse coding. The measures were short, study-specific indicators inspired by relevant literature. They were not direct adaptations of complete published instruments and are not presented as established psychometric measures.

6.3.1 AI use

AI-use frequency was recorded on an ordinal scale from never to daily or almost daily. Respondents also indicated stages of use, types of AI tools, perceived competence, and the declared share of AI in the final output. AI final share was measured from 1 (none or almost none) to 5 (very large). Analyses of H1-H3 used creators who reported using AI because personal control and AI share have their clearest meaning in this subgroup.

6.3.2 Perceived authorship

The primary authorship score averaged three items addressing whether an AI-assisted result could remain one's work and whether greater AI involvement or predominantly generated output made the work feel less one's own. Two negatively phrased items were reverse coded. A fourth item about being the author when making key creative decisions was treated as an auxiliary indicator because it more directly foregrounded the proposed mechanism. The three-item score had low internal consistency, Cronbach's $\alpha = .519$ and $\omega = .745$.

6.3.3 Perceived authenticity

The three-item authenticity score addressed whether AI-assisted work could remain authentic when it expressed human intention, whether it could preserve personal style, and whether AI use weakened authenticity even when the idea came from a person. The negative item was reverse coded. A separate item concerning a recognizable "typical AI aesthetic" was analyzed as an auxiliary concern rather than included in the score. The primary score had low internal consistency, $\alpha = .521$ and $\omega = .755$.

6.3.4 Perceived control

The four-item control indicator assessed whether the respondent felt they directed the process, could consciously select, revise, and reject AI proposals, experienced AI as taking the work in an unwanted direction, and regarded the final form as dependent mainly on their decisions. One item was reverse coded. Internal consistency was low, $\alpha = .452$ and $\omega = .677$. The negative corrected item-total correlation for the reverse-coded item indicated that the score should be interpreted as a short composite of related control judgments rather than as a homogeneous psychometric scale.

6.3.5 Creative identity threat

The four-item identity-threat indicator concerned anxiety about the value of human creative work, the distinctiveness of creative competence, the meaning of being a creator, and viewing AI as an extension of one's toolkit rather than a threat. The last item was reverse coded. Reliability was limited, $\alpha = .646$ and $\omega = .793$. The reverse-coded item showed little corrected item-total association, so conclusions involving the composite remain cautious.

6.3.6 Expressive orientation toward art

Three items assessed the view that art should express human personal experience, intention, and sensibility, and that personal meaning can matter more than technical execution. A fourth reverse-worded item about the possible value of a work without the author's personal experience was treated as auxiliary. The three-item score had acceptable internal consistency, $\alpha = .727$ and $\omega = .850$.

6.3.7 AI as tool and AI usefulness

The two-item AI-as-tool index assessed whether AI was viewed as a tool or assistant rather than an independent creator and whether human selection and evaluation remained most important. Its α was .828, although any two-item index has limited psychometric breadth and should be interpreted cautiously.

The three-item AI-usefulness score addressed acceleration from idea to first version, exploration of variants the creator might not otherwise consider, and expansion of human creative possibilities when a person set the direction. Reliability was good, $\alpha = .850$ and $\omega = .910$.

6.3.8 Vignettes

Each vignette was followed by three ratings: whether the final work could be regarded as the person's work, whether it could authentically express the person's intention, and whether the person retained control. The mean of the three items formed the controlled-use and final-generator indices. Both indices showed good internal consistency, $\alpha = .892$; ω was .933 for each. One respondent had incomplete controlled-vignette data, resulting in 428 complete paired cases for the comparison.

Table 2. Reliability of study indicators

Indicator	Items	Complete n	Cronbach's alpha	Omega	Interpretation
Authorship	3	429	0.519	0.745	Low reliability
Authenticity	3	429	0.521	0.755	Low reliability
Control	4	429	0.452	0.677	Low reliability
Creative identity threat	4	429	0.646	0.793	Limited reliability
Expressive orientation	3	429	0.727	0.850	Acceptable reliability
AI as tool	2	429	0.828		Short two-item index; interpret cautiously
AI usefulness	3	429	0.850	0.910	Good reliability
Controlled-use vignette	3	428	0.892	0.933	Good reliability
Final-generator vignette	3	429	0.892	0.933	Good reliability

6.4 Procedure

Before the main study, the questionnaire underwent a small pilot focused on completion time, wording, vignette clarity, conditional logic, and data export. Pilot data were excluded from the main dataset. After minor wording adjustments, the questionnaire was frozen as `questionnaire_v1.0` before preregistration and public recruitment.

The Polish-language questionnaire was administered on the Ankieto platform. Its frozen public version contained 46 system-numbered elements, including the consent question, distributed across 13 screens. Participants received a general description of the study and were required to confirm that they were adults and consented to participate; refusal of consent ended participation. Participation was voluntary, and respondents could stop at any time. The questionnaire contained demographic and creative-profile items, questions about AI use, attitudinal indicators, two vignettes, and three optional final open-ended questions. Open-ended responses were not used in the quantitative analyses reported here.

All 429 completed and consenting records in the final export were retained in the primary analysis. IP addresses were used only in a private technical check to flag potential duplicate submissions; they were excluded from the analytic dataset and public materials, and a shared IP address did not trigger automatic exclusion. Very short or very long completion times likewise did not trigger automatic removal from the primary analysis. Instead, the preregistered results were calculated on the full analytic sample and compared with sensitivity analyses excluding responses below two minutes, below three minutes, and outside a three-minute to 120-minute range. This approach recognized that very short responses may indicate insufficient attention, while very long durations may reflect interruptions rather than low-quality responding.

6.5 Preregistration, data handling, and ethics

The study was preregistered on the Open Science Framework before the first recorded main-study response: <https://doi.org/10.17605/OSF.IO/GSWN3>. H1-H4 were designated confirmatory, while H5-H6 and additional group comparisons were designated exploratory.

The analysis used a de-identified quantitative dataset. Direct technical identifiers, exact survey-platform identifiers, browser and system information, exact timestamps, and full open-ended responses were removed or excluded from the analytic file. The raw survey export was not used for substantive statistical analysis after the de-identification workflow and is not publicly shared.

No formal ethics committee approval or formal exemption was obtained before data collection.

Participants were adults and received information about the general purpose of the study before providing informed consent. Participation was voluntary, and respondents could stop completing the survey. Names and contact details were not requested. Platform-generated IP addresses were used solely for a private duplicate check, removed from the analytic dataset, and not publicly shared. The analyses used a de-identified quantitative dataset, and the raw survey export is not public. The survey did not intentionally

request medical, clinical, or other sensitive information. These safeguards do not constitute prospective ethics approval or exemption.

6.6 Data availability and open-science statement

The preregistered scoring was retained despite low reliability in several short indicators, and all post hoc psychometric and model checks are labelled as exploratory robustness analyses rather than replacements for H1-H4. The raw survey export is not publicly available because it contains platform-generated technical metadata, including IP addresses. The analyses were conducted on a de-identified quantitative dataset from which direct technical identifiers, exact timestamps, browser and system information, and full open-ended responses were removed. The de-identified quantitative dataset and analysis scripts are not publicly available at the time of publication and may be shared through an appropriate research repository following an additional privacy review. The preregistration materials, including the codebook, are publicly available on OSF: <https://doi.org/10.17605/OSF.IO/GSWN3>. The full questionnaire, scoring key, documented procedures, participant-flow limits, and preregistration concordance are provided in Appendix A.

6.7 Use of AI-assisted tools and author responsibility

AI-assisted tools were used only for language editing, translation, and manuscript assembly. All research, analytical, interpretive, and editorial decisions were made and verified by the author, who takes full responsibility for the manuscript. AI-assisted output did not determine hypothesis classifications, exclusions, model selection, or interpretation.

6.8 Statistical analysis

Analyses were conducted in Python using pandas, SciPy, and statsmodels. Reliability was assessed with Cronbach's alpha, a one-factor estimate of McDonald's omega for indicators with at least three items, and corrected item-total correlations. Descriptive statistics were calculated for the full sample and relevant subgroups.

For H1 and H2, the primary tests were Pearson correlations in the creator-AI-user subgroup; Spearman correlations and simple ordinary least squares regressions were retained as supporting analyses. H3 was tested using ordinary least squares regression predicting authenticity from mean-centered AI final share, mean-centered control, and their interaction. H4 was examined in the full sample using Pearson correlations and regression, with identity threat as the primary outcome and the auxiliary typical-AI-aesthetic concern as a related outcome.

The two vignette indices and each dimension were compared using paired-samples t tests. Wilcoxon signed-rank tests were calculated as robustness checks, and Cohen's d_z quantified the standardized within-person difference. Exploratory analyses examined the association of AI usefulness with identity threat and the ordinal association of AI-use frequency with the AI-as-tool index. Tests were two-sided at

alpha = .05. Because H4 allowed more than one theoretically related indicator and no familywise correction was applied, its interpretation is explicitly cautious. Post hoc robustness checks using HC3 standard errors and bootstrap confidence intervals did not change the classification of H2-H4.

7 Results

7.1 Descriptive statistics and reliability

All 429 records had complete values for the primary indicator scores, except that one respondent lacked a complete controlled-vignette index. Across the full sample, mean authorship was 2.61 (SD = 0.91), authenticity 2.88 (SD = 0.86), control 3.26 (SD = 0.71), identity threat 3.23 (SD = 0.84), expressive orientation 3.65 (SD = 0.90), AI as tool 3.69 (SD = 1.04), and AI usefulness 3.47 (SD = 0.99). These means summarize a heterogeneous sample and are not interpreted as population estimates.

Reliability varied substantially. Expressive orientation was acceptable, AI usefulness was good, and the two vignette indices were good. In contrast, the authorship, authenticity, and control composites had low alpha values. Their omega estimates were higher, but the item-level pattern showed that some items contributed weakly or negatively after reverse coding. Identity threat also included a weak reverse-coded item. These results do not justify deleting items after viewing the data and then presenting a more favorable post hoc scale. The preregistered scores were retained, and the limitations are carried into the interpretation of H1-H4.

7.2 Preregistered confirmatory analyses

7.2.1 H1: AI final share and perceived authorship

H1 predicted a negative association between declared AI final share and authorship among creators who used AI. The observed Pearson correlation was positive and did not reach the .05 threshold, $r = .141$, 95% CI $[-.012, .288]$, $p = .071$, $N = 164$. Spearman's rho was $.131$, $p = .094$. In a simple regression, AI final share predicted a small positive authorship difference, $b = .109$, $SE = .060$, $p = .071$, $R^2 = .020$.

H1 was not supported. The direction was opposite to the preregistered prediction. This result should not be re-described as evidence that greater AI involvement increases authorship, because the confidence interval included zero and the measure of authorship had low reliability. It does, however, argue against a simple assumption that declared AI share necessarily reduces felt authorship in this subgroup.

7.2.2 H2: perceived control and perceived authorship

H2 predicted a positive association between control and authorship among creators who used AI. The Pearson correlation supported this prediction, $r = .234$, 95% CI $[.084, .374]$, $p = .003$, $N = 164$. In the corresponding simple regression, $b = .275$, $SE = .090$, $p = .003$, $R^2 = .055$. The full-sample association, treated as auxiliary because non-users may answer personal control items more hypothetically, was smaller but positive, $r = .142$, $p = .003$, $N = 429$.

H2 was supported. The effect was modest, and the low reliability of both composites means that its magnitude should not be treated as a precise estimate. The result is nonetheless consistent with the proposed importance of retaining direction over the process.

7.2.3 H3: control as a moderator

H3 predicted that control would weaken a negative association between AI final share and authenticity. The model for creator AI users included centered AI share, centered control, and their interaction. The overall model explained 17.5% of variance in the authenticity indicator. Control had a positive main association in the model, $b = .431$, $SE = .079$, $p < .001$. The interaction was not significant, $b = .043$, $SE = .073$, 95% CI $[-.102, .187]$, $p = .559$, $N = 164$. H3 was therefore not supported in the main analysis.

The non-significant interaction is important for interpretation. The data do not establish that control statistically buffers a harmful effect of AI share on authenticity. The positive main association of control with authenticity is compatible with the broader theory, but it does not substitute for the predicted interaction.

7.2.4 H4: expressive orientation and concern about AI

H4 predicted that stronger expressive orientation would be associated with greater concern about authenticity and/or creative identity threat. In the full sample, expressive orientation correlated positively with identity threat, $r = .341$, 95% CI $[.254, .422]$, $p < .001$, $N = 429$. It was also associated with the auxiliary item measuring concern that a typical AI aesthetic makes work seem less authentic, $r = .504$, 95% CI $[.430, .572]$, $p < .001$. By contrast, expressive orientation was not associated with the general three-item authenticity score, $r = .011$, $p = .827$.

H4 was cautiously supported. The identity-threat association was in the predicted direction and stable, while the stronger auxiliary association indicates sensitivity to perceived AI aesthetics. The null result for general authenticity and the limited reliability of identity threat argue against a broad claim that expressive values uniformly reduce perceived authenticity of AI-assisted work.

Table 3. Preregistered confirmatory analyses

Hypothesis	Sample	Test	Estimate	95% CI low	95% CI high	p	Interpretation
H1: AI final share with authorship	Creator AI users, N = 164	Pearson r	0.141	-0.012	0.288	0.071	Not supported; direction opposite to prediction
H2: control with authorship	Creator AI users, N = 164	Pearson r	0.234	0.084	0.374	0.003	Supported with reliability caveat
H3: AI share x control predicting authenticity	Creator AI users, N = 164	OLS interaction b	0.043	-0.102	0.187	0.559	Not supported; model R ² = 0.175
H4: expressive orientation with identity threat	Full sample, N = 429	Pearson r	0.341	0.254	0.422	< .001	Cautiously supported
H4 auxiliary: expressive orientation with concern about typical AI aesthetics	Full sample, N = 429	Pearson r	0.504	0.430	0.572	< .001	Consistent with H4; auxiliary item

7.3 Vignette analyses

The vignette comparison produced the clearest pattern in the study. The combined rating for human-controlled AI use was $M = 3.501$ ($SD = 1.081$), compared with $M = 2.665$ ($SD = 1.132$) for the final-generator scenario. The mean paired difference was 0.836, 95% CI [0.715, 0.956], $t(427) = 13.628$, $p < .001$, Cohen's $d_z = 0.659$. The Wilcoxon signed-rank test also indicated a difference, $W = 5581.5$, $p < .001$.

The pattern held across all three dimensions. The controlled scenario was rated .860 points higher on authorship, $t(427) = 12.625$, $p < .001$, $d_z = .610$; .801 points higher on authenticity, $t(427) = 11.575$, $p < .001$, $d_z = .559$; and .846 points higher on control, $t(427) = 12.360$, $p < .001$, $d_z = .597$. These convergent differences make the vignette result the strongest and most stable finding in the dataset.

The comparison should still be interpreted as a judgment task, not as an experiment demonstrating what happens to creators who adopt a particular workflow. The scenarios differed in several intentionally bundled features, including initial intention, selection of variants, degree of editing, and acceptance of the first output. The result identifies the package of human direction as consequential for evaluation; it does not isolate one causal component.




Fig. 1 Mean evaluation of the human-controlled and final-generator vignettes. Error bars show one standard deviation




Fig. 2 Mean ratings of authorship, authenticity, and control in the two vignettes. Error bars show one standard deviation

Table 4. Within-person vignette comparisons

Dimension	Paired n	Controlled M	Controlled SD	Generator M	Generator SD	Mean difference	t	p	Cohen's dz
Authorship	428	3.430	1.208	2.570	1.235	0.860	12.625	< .001	0.610
Authenticity	428	3.507	1.172	2.706	1.253	0.801	11.575	< .001	0.559
Control	428	3.565	1.196	2.720	1.255	0.846	12.360	< .001	0.597
Combined index	428	3.501	1.081	2.665	1.132	0.836	13.628	< .001	0.659

7.4 Exploratory analyses

H5 concerned the possible coexistence of perceived usefulness and creative identity threat. AI usefulness was not correlated with identity threat in the full sample, $r = -.006$, $p = .901$, $N = 429$; Spearman's rho was $.032$, $p = .510$. Using a descriptive threshold of at least 4 on both indices, 10.26% of respondents scored high on both. This subgroup shows that positive and negative evaluations can coexist for some participants, but the absence of a general association does not support a single-dimensional construct of ambivalence.

H6 received limited exploratory support. More frequent AI use was weakly associated with seeing AI as a tool or assistant, Spearman's rho = $.154$, $p = .001$, $N = 429$. The association is small, and frequency groups were uneven. It is compatible with familiarization or self-selection, but the cross-sectional design cannot distinguish whether use changes views or whether people who already view AI as a tool use it more often.

Exploratory group comparisons showed that AI users and more active creator AI users often reported higher authorship, authenticity, control, and AI usefulness than their comparison groups. These patterns may help explain why H1 did not show the predicted negative direction: experienced users may be a selected group that has developed workflows in which AI involvement and human direction coexist. Because these comparisons were not the primary confirmatory tests and involved multiple outcomes, they are treated as hypothesis-generating.

7.5 Completion-time sensitivity

Time filters did not replace the preregistered primary analysis. After excluding responses below two minutes, below three minutes, or outside three to 120 minutes, H2, H4, and the vignette difference remained directionally stable. H1 became statistically significant after short responses were removed, but the association remained positive and therefore opposite to the prediction. H3 remained unsupported; its interaction approached but did not cross the conventional threshold under several filters (p values $.052$ -. $.069$). This sensitivity pattern reinforces the decision to report H1 as not supported and H3 as not supported in the main analysis, with H3 noted as filter-sensitive rather than reclassified.

8 Discussion

The study examined whether evaluations of AI-assisted creative practice are better understood through the amount of AI involvement or through the structure of human control. The preregistered results do not provide a simple confirmation of the proposed model. H1 and H3 were not supported. H2 showed a modest positive association between perceived control and authorship among creators who used AI. H4 was cautiously supported through associations of expressive orientation with identity threat and concern about a typical AI aesthetic, but not with the general authenticity indicator. The clearest evidence came

from the vignette comparison, in which respondents consistently preferred a human-directed process over acceptance of a near-final generated output.

This mixed pattern does not establish meaningful human control as a causal mechanism, validated model, or general standard. It supports moving beyond a single AI quantity or label and treating authorship and authenticity as responses that may also depend on how a process is organized.

8.1 Why AI involvement alone was not sufficient to explain authorship

H1 predicted that greater declared AI share would be associated with lower authorship among creator AI users. The observed association was small, non-significant, and positive. The result does not show that greater AI involvement increases authorship, but it directly fails to support the preregistered “more AI, less authorship” prediction.

AI final share may be too coarse to capture the contribution that matters. It records perceived quantity, not whether the system supplied a peripheral component or a decision central to the work. A creator can incorporate extensive generated material while maintaining a strong concept, curating alternatives, and transforming the final artifact. A comparatively small generated contribution can instead determine the narrative, composition, or style. This qualitative difference is consistent with attribution research emphasizing the significance of contributions rather than participation alone (He et al. 2025).

This distinction also clarifies why the result need not replicate manipulated-assistance vignette studies. Formosa et al. (2025) varied how much lightly edited assistant-written text remained in a fixed scenario and found lower human authorship at higher assistance levels; H1 used creators' retrospective estimates of AI share across heterogeneous real practices. The two approaches capture related but non-equivalent aspects of contribution.

8.2 Human control as an interpretive lens

H2 provides direct but modest support for control. Among creators using AI, higher perceived control was associated with higher perceived authorship. The result aligns with a situated sense-of-agency account: work is more likely to feel one's own when intention, choice, and outcome remain connected (Tapal et al. 2017). It also supports examining control as a feature of creative practice rather than treating authorship as a simple residue left after subtracting AI involvement.

The low reliability of the control and authorship indicators substantially limits this inference. Control may include separable experiences of goal setting, predictability, veto power, editability, and final authority. The present short indicator did not measure direction, selection, revision, and responsibility as distinct validated dimensions. The process account should therefore be read as a theoretically organized interpretation of converging evidence, not as a latent structure confirmed by these data.

H3 is an important boundary. The preregistered claim was not merely that control correlates with positive evaluations; it was that control would weaken a negative relation between AI share and authenticity. The

interaction was not significant. A positive main association of control with authenticity cannot replace the unsupported interaction. Control may have an additive relation with authenticity, the measures may be inadequate, or the moderation may be absent. These alternatives require new data and should not be resolved post hoc.

8.3 What the vignette contrast reveals about direction, selection, and revision

The vignette contrast provides the strongest evidence for process-sensitive evaluation. Respondents rated the human-directed scenario higher on authorship, authenticity, and control, with a combined within-person effect of $d_z = .659$. The scenarios shared the broad fact of AI use but differed in how the process was organized. The controlled scenario included an original idea, generation of variants, human selection, substantial editing, and final compositional authority. The final-generator scenario included a short brief, acceptance of the first result, and only minor changes.

This result demonstrates that respondents attend to more than AI presence. It does not demonstrate which component caused the difference. Direction, selection, and revision were intentionally bundled, and responsibility was implied rather than independently manipulated. Demand characteristics and the greater narrative detail of the controlled scenario may also contribute. Nonetheless, the convergence across three outcomes shows that the bundled process distinction was legible and consequential.

Experimental-vignette guidance emphasizes realistic scenarios and controlled manipulations that support both internal and external validity (Aguinis and Bradley 2014). The present within-person contrast provides a legible process package, but it cannot estimate the independent contribution of each bundled feature.

The result can be interpreted as evidence supporting a process-sensitive account of AI-assisted creative practice. A workflow appears more human-directed when the person contributes an intention, evaluates alternatives, materially intervenes, and determines the final form. Future factorial experiments should separate these features and test whether each has an independent or interactive effect.

8.4 Creative-tool design should be process-sensitive, not binary

Binary labels can support transparency but are poor design specifications. An interface cannot preserve agency simply by placing a person somewhere “in the loop.” If the person sees only one polished output, cannot inspect alternatives, and has no practical means to revise central features, a confirmation button may provide nominal rather than meaningful control.

The findings do not determine one correct interface. Creative fields differ in what counts as a central decision, and users differ in goals and expertise. The empirical contribution is narrower: people evaluated two AI-assisted workflows differently when human direction and transformation differed. This supports treating workflow affordances as a design concern alongside output quality and productivity.

8.5 Implications for attribution and disclosure

Disclosure that AI was used may be necessary, but it is not sufficient for attribution. A binary declaration can group together ideation support, component generation, extensive co-editing, and near-complete output generation. More informative disclosure can describe the role of AI in plain process terms: whether it supported brainstorming, generated variants, produced specific components, or supplied a near-final artifact; and whether the human selected and substantially revised the output.

Process descriptions should remain concise. A brief account of direction, selection, revision, and responsibility may be more informative than exact percentages across unlike media. Attribution should also distinguish credit from accountability: claimed authorship and assigned responsibility should reflect actual authority to inspect, alter, and approve the work.

8.6 Implications for creative professionals and education

For creative professionals, consequential control is partly an organizational resource. Deadlines, client demands, and managerial policies can shape whether AI remains optional assistance or becomes de facto automation. Workflows can clarify who decides to use AI, who reviews outputs, whether creators can reject unsuitable generations, and who bears accountability.

The exploratory results caution against dividing creators into pro-AI and anti-AI camps. AI usefulness and identity threat were essentially unrelated in the full sample, while a minority scored high on both. Frequent use was only weakly associated with seeing AI as a tool. Creators may value efficiency and exploration while worrying about devaluation of labor or homogenization. Professional practice should accommodate this coexistence rather than treating concern as lack of competence.

For education, relevant capabilities include formulating intentions, evaluating alternatives, revising generated material, documenting decisions, and explaining one's contribution. Process records and reflective justification can connect craft with judgment while supporting critical engagement with identity, agency, and ethics in AI-assisted art education (Lee et al. 2026).

8.7 Implications for interfaces and workflow design

The four-part interpretive lens suggests bounded interface questions: can users preserve goals across iterations, compare alternatives, revise central features, reject outputs, and review a usable record of generated components and human approvals? These questions are consistent with control-by-design accounts emphasizing feedback and continued human learning (Te'eni et al. 2026) and with human-AI interaction guidance across stages of use (Amershi et al. 2019). More generated options should not be equated with more control when users cannot navigate or revise them.

9 Limitations

Several limitations constrain the conclusions. First, the sample was non-probability and recruited online. It cannot estimate the prevalence of attitudes among creators or the general population. Creative fields were uneven, and some categories were too small for meaningful comparison. The relatively large proportion of respondents aged 35 and older may also shape exposure to AI and professional identity.

Second, the design was cross-sectional and primarily self-report. Correlations cannot establish whether control increases authorship, whether authorship encourages perceptions of control, or whether both reflect experience, confidence, or preferred workflow. Declared AI share is subjective and may be interpreted differently across media. The study did not inspect actual artifacts or production histories.

Third, measurement quality was mixed. The authorship, authenticity, and control indicators had low alpha values, and identity threat had limited reliability. Reverse-worded items contributed to some problems. These indicators are not validated scales, their absolute scores should not be overinterpreted, and measurement error may attenuate associations or destabilize interaction tests. The higher omega estimates do not remove the item-level concerns.

Fourth, the vignette comparison was not a randomized factorial experiment. Both scenarios were shown in the questionnaire, and the contrast bundled intention, alternatives, editing, and final authority. Order or demand characteristics may have amplified the difference. The graphic-design example may have been more concrete for visual creators than for those working in music, writing, or performance. The vignettes also differed in protagonist gender and task context: the human-directed scenario described a female creator developing her own idea, whereas the final-generator scenario described a male creator responding to a commissioned brief. These differences were not independently controlled and may have contributed to the observed contrast.

Fifth, completion times varied widely. Sensitivity analyses reduced but did not eliminate this concern. H2, H4, and the vignette contrast were directionally robust. H1 and H3 varied under time filters, but the sensitivity results did not justify changing their preregistered classifications.

Sixth, the analysis included multiple related outcomes and exploratory comparisons without a global multiplicity correction. H4's "and/or" formulation allowed more than one indicator. This flexibility, the null general-authenticity association, and limited reliability are why H4 is described as cautiously supported.

Finally, no formal ethics committee approval or formal exemption was obtained before data collection. Participants were adults, names and contact details were not requested, platform-generated IP addresses were used only for a private duplicate check and removed from the analytic dataset, and informed consent was obtained. Participation was voluntary, and respondents could stop completing the survey. The analyses used a de-identified quantitative dataset, and the raw survey export is not public. These

safeguards do not substitute for prospective institutional review where required and do not constitute retrospective approval or exemption.

10 Future research

The next step is not to declare a standard but to test the framework more directly. First, measurement development should distinguish direction, selection, revision, and responsibility using cognitive interviews, larger item pools, preregistered psychometric analyses, and independent validation samples. Perceived control should be separated from technical confidence, predictability, and general attitudes toward AI.

Second, experiments should independently manipulate idea origin, number and diversity of alternatives, degree of revision, final veto power, and responsibility. Counterbalanced, domain-neutral, and field-specific vignettes can identify which components shape authorship and authenticity judgments. Behavioral studies can examine actual interaction histories rather than relying only on self-report.

Third, longitudinal research can distinguish self-selection from adaptation. Creators who adopt AI may develop practices that preserve agency, while those who feel displaced may stop using it. Repeated measures across tool adoption would show whether meaningful control changes and whether changes predict professional identity, satisfaction, or continued use.

Finally, human control should be studied alongside other conditions of AI-assisted production, including data provenance, consent, compensation, bias, accessibility, labor conditions, and environmental cost. A process can be human-directed yet remain problematic on other grounds. The value of the process lens lies in adding a neglected dimension, not replacing broader ethical assessment.

11 Conclusion

The preregistered evidence does not support a simple rule that more AI involvement produces less human authorship. H1 was not supported and ran opposite to its predicted direction; H3's moderation was also unsupported. Perceived control showed a modest positive association with authorship, and respondents clearly distinguished a human-directed workflow from acceptance of a near-final generated output. Expressive orientation was associated with identity threat, although measurement and multiplicity limitations require caution.

Together, the findings support a process-sensitive account of AI-assisted creative practice, but they do not establish a definitive standard or causal mechanism. Declared AI quantity was insufficient to explain perceived authorship in the predicted way, whereas perceived control showed a modest association and the bundled process contrast produced the clearest result. Direction, selection, revision, and responsibility offer a useful vocabulary for interpreting that pattern, not a validated scale. Creative tools, workflows, education, and professional practice can treat consequential human decision-making as a testable design

concern while avoiding the assumption that an AI label or percentage settles questions of authorship and authenticity.

References

- Aguinis H, Bradley KJ (2014) Best practice recommendations for designing and implementing experimental vignette methodology studies. *Organizational Research Methods* 17(4):351-371. <https://doi.org/10.1177/1094428114547952>
- Amershi S, Weld D, Vorvoreanu M, Fournery A, Nushi B, Collisson P, Suh J, Iqbal S, Bennett PN, Inkpen K, Teevan J, Kikin-Gil R, Horvitz E (2019) Guidelines for human-AI interaction. In: *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*. ACM, pp 1-13. <https://doi.org/10.1145/3290605.3300233>
- Cavalcante Siebert L, Lupetti ML, Aizenberg E, Beckers N, Zgonnikov A, Veluwenkamp H, Abbink D, Giaccardi E, Houben GJ, Jonker CM, van den Hoven J, Forster D, Lagendijk RL (2023) Meaningful human control: Actionable properties for AI system development. *AI and Ethics* 3:241-255. <https://doi.org/10.1007/s43681-022-00167-3>
- Cherry E, Latulipe C (2014) Quantifying the creativity support of digital tools through the Creativity Support Index. *ACM Transactions on Computer-Human Interaction* 21(4):1-25. <https://doi.org/10.1145/2617588>
- Cox SR, Djernæs HB, van Berkel N (2025) Beyond productivity: Rethinking the impact of creativity support tools. In: *Proceedings of the 2025 Conference on Creativity and Cognition*. ACM, pp 735-749. <https://doi.org/10.1145/3698061.3726924>
- Draxler F, Werner A, Lehmann F, Hoppe M, Schmidt A, Buschek D, Welsch R (2024) The AI Ghostwriter Effect: When users do not perceive ownership of AI-generated text but self-declare as authors. *ACM Transactions on Computer-Human Interaction* 31(2):1-40. <https://doi.org/10.1145/3637875>
- Formosa P, Bankins S, Matulionyte R, Ghasemi O (2025) Can ChatGPT be an author? Generative AI creative writing assistance and perceptions of authorship, creatorship, responsibility, and disclosure. *AI & SOCIETY* 40:3405-3417. <https://doi.org/10.1007/s00146-024-02081-0>
- Haidemariam T (2026) After generative AI: Authorship, labour, and cultural governance. *AI & SOCIETY*. <https://doi.org/10.1007/s00146-026-03189-1>
- He J, Houde S, Weisz JD (2025) Which contributions deserve credit? Perceptions of attribution in human-AI co-creation. In: *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. ACM, pp 1-18. <https://doi.org/10.1145/3706598.3713522>
- Hwang AHC, Liao QV, Blodgett SL, Olteanu A, Trischler A (2025) “It was 80% me, 20% AI”: Seeking authenticity in co-writing with large language models. *Proceedings of the ACM on Human-Computer Interaction* 9(2):1-41. <https://doi.org/10.1145/3711020>
- Jiang HH, Taylor J, Agnew W (2026) How professional visual artists are negotiating generative AI in the workplace. In: *Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems*, pp 1-6. <https://doi.org/10.1145/3772363.3799003>
- Kieslich K, Lünich M, Marcinkowski F (2021) The Threats of Artificial Intelligence Scale (TAI): Development, measurement and test over three application domains. *International Journal of Social Robotics* 13(7):1563-1577. <https://doi.org/10.1007/s12369-020-00734-w>
- Lee H, Yu W, Lee J (2026) Negotiating creator identity: Agency and ethical awareness in AI-assisted art education. *AI & SOCIETY*. <https://doi.org/10.1007/s00146-026-03187-3>
- Lovato J, Zimmerman JW, Smith I, Dodds P, Karson JL (2024) Foregrounding artist opinions: A survey study on transparency, ownership, and fairness in AI generative art. *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society* 7(1):905-916. <https://doi.org/10.1609/aies.v7i1.31691>
- Polimetla T, Gero KI, Glassman EL (2026) A paradigm for creative ownership. In: *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems*. ACM, Article 148, pp 1-16. <https://doi.org/10.1145/3772318.3791316>

- Salas Espasa D, Camacho M (2025) From aura to semi-aura: Reframing authenticity in AI-generated art - a systematic literature review. *AI & SOCIETY* 40(8):6727-6759. <https://doi.org/10.1007/s00146-025-02361-3>
- Santoni de Sio F, van den Hoven J (2018) Meaningful human control over autonomous systems: A philosophical account. *Frontiers in Robotics and AI* 5:15. <https://doi.org/10.3389/frobt.2018.00015>
- Tapal A, Oren E, Dar R, Eitam B (2017) The Sense of Agency Scale: A measure of consciously perceived control over one's mind, body, and the immediate environment. *Frontiers in Psychology* 8:1552. <https://doi.org/10.3389/fpsyg.2017.01552>
- Te'eni D, Yahav I, Schwartz D (2026) What it takes to control AI by design: Human learning. *AI & SOCIETY* 41:237-250. <https://doi.org/10.1007/s00146-025-02401-y>
- Tsao J, Liang CX, Nogues C, Wong A (2026) Perceptions and integration of generative artificial intelligence in creative practices and industries: A scoping review and conceptual model. *AI & SOCIETY* 41:2259-2278. <https://doi.org/10.1007/s00146-025-02667-2>
- Weaver A (2026) Sisters, not twins: Exploring artistic control and anthropomorphism through composing with a bespoke generative AI. *AI & SOCIETY* 41:1609-1621. <https://doi.org/10.1007/s00146-025-02291-0>